

Task 3: Sales Prediction

3Q) Sales prediction means predicting how much of a product people will buy based on factors such as the amount you spend to advertise your product, the segment of people you advertise for, or the platform you are advertising on about your product. Typically, a product and service-based business always need their Data Scientist to predict their future sales with every step they take to manipulate the cost of advertising their product. So let's start the task of sales prediction with machine learning using Python.

```
In [1]: 1 # importing all required packages
        2
        3 import numpy as np
        4 import pandas as pd
```

```
In [2]: 1 # reading the dataset
        2
        3 df=pd.read_csv('Advertising.csv')
        4 df.head()
```

```
Out[2]:
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9

```
In [3]: 1 df.shape
```

```
Out[3]: (200, 5)
```

```
In [4]: 1 df.dtypes # identifying the datatypes of each column in the dataset
```

```
Out[4]: Unnamed: 0      int64
TV          float64
Radio       float64
Newspaper   float64
Sales       float64
dtype: object
```

```
In [5]: 1 df.isnull().sum() # checking the number of null values in the data, in
```

```
Out[5]: Unnamed: 0      0
TV          0
Radio       0
Newspaper   0
Sales       0
dtype: int64
```

```
In [6]: 1 df.drop('Unnamed: 0',axis=1,inplace=True) # drop id column to reduce cl
```

```
In [7]: 1 df.head()
```

```
Out[7]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

```
In [8]: 1 df.shape
```

```
Out[8]: (200, 4)
```

```
In [9]: 1 df.dtypes
```

```
Out[9]: TV          float64
Radio          float64
Newspaper      float64
Sales          float64
dtype: object
```

```
In [10]: 1 X=df.drop('Sales',axis=1) # seperating input columns from output column
2 X.head() # here, input column consists of TV, Radio, Newspaper
```

```
Out[10]:
```

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4

```
In [11]: 1 y=df['Sales'] # here, the output column is sales
2 y.head()
```

```
Out[11]: 0    22.1
1    10.4
2     9.3
3    18.5
4    12.9
Name: Sales, dtype: float64
```

```
In [12]: 1 # Splitting data into train and test data
2
3 from sklearn.model_selection import train_test_split
4 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random
```

```
In [13]: 1 X_train.shape,y_train.shape
```

```
Out[13]: ((140, 3), (140,))
```

```
In [14]: 1 X_test.shape,y_test.shape
```

```
Out[14]: ((60, 3), (60,))
```

A) Linear Regression

```
In [15]: 1 # using a linear regression model
2
3 from sklearn.linear_model import LinearRegression
4 LR=LinearRegression()
5 LR.fit(X_train,y_train) # fit the train data into the model
```

```
Out[15]: LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [16]: 1 y_train_pred_LR=LR.predict(X_train) # use the fit model to predict the
2 y_train_pred_LR
```

```
Out[16]: array([17.39149783, 15.19196153, 11.41650701, 11.20610472, 16.39256165,
        6.90577778, 21.17740606,  6.10528574,  9.66662607, 11.67308587,
        9.0704377 ,  6.30582199, 14.7930121 , 17.42999512, 16.14599956,
        16.12273906, 14.9308629 , 19.61593142, 13.70741553, 21.09067507,
        13.09101877, 13.79551693,  8.87971636, 16.97348947,  8.22902448,
        15.34026923, 13.95862675, 23.23157581, 12.67992504, 23.10768546,
        6.80057243, 18.81832259, 23.69669553, 18.39890879, 16.97890645,
        16.44055305, 12.41657918, 11.94072527, 16.7732918 , 14.5933997 ,
        13.22910727,  7.49691601, 19.30121038,  9.33105452, 19.3603766 ,
        10.16463427,  6.94369039, 16.52918217, 13.53571009, 14.80225851,
        11.13948107, 20.85632272, 24.02985438, 18.427486 , 17.84570024,
        15.11333638, 17.21352856,  9.2345359 , 17.29735156, 19.22919752,
        16.29255016,  3.54203145,  5.24055709, 15.86268553, 15.12384811,
        16.31596188, 18.34249185, 23.4941462 , 14.34540589, 20.48784226,
        17.29893239, 21.67652708, 10.30595955, 15.25363154, 13.68052604,
        9.90613461, 18.01712723, 18.43772847, 20.3652561 , 19.37970879,
        10.18390434, 15.52162749, 20.86386386,  8.24145294, 11.54953301,
        7.25535492, 11.99074605, 14.68055076,  4.4950967 , 14.83591282,
        19.81147067, 13.60063366, 17.42635854,  3.5132451 ,  8.56886762,
        17.93522755, 14.44803773, 15.26734959, 23.38305647, 10.59905655,
        24.94180909, 12.47872281, 12.99022295, 24.44392073, 13.19055975,
        20.1322331 ,  6.91136723, 12.36791752,  7.89158996, 22.14077381,
        16.21139454, 12.18638677, 14.34672625, 12.65039181, 14.12757003,
        4.50133295,  8.02296589, 16.15950779,  9.73534671, 18.45288885,
        12.81198432, 20.82621122,  9.96179489, 10.6914528 , 10.04931249,
        14.53862883, 17.29195706, 11.86800298, 16.11235485, 17.1038322 ,
        8.20867785, 14.70885801, 18.21928267, 18.10546473, 10.61464308,
        6.20700241, 18.57459979, 19.38285038, 12.11917215, 17.21444781])
```

```
In [17]: 1 # compare the y_train data and y_train_pred_LR data to check the perfor
2
3 from sklearn.metrics import r2_score,mean_absolute_error
4 R2=r2_score(y_train,y_train_pred_LR)
5 MSE=mean_absolute_error(y_train,y_train_pred_LR)
6 RMSE=np.sqrt(MSE)
7 print('R-Square: ',R2)
8 print('Mean Square Error: ',MSE)
9 print('Root Mean Square Error: ',RMSE)
```

R-Square: 0.9055159502227753
Mean Square Error: 1.158150294807253
Root Mean Square Error: 1.076173914758787

```
In [18]: 1 y_test_pred_LR=LR.predict(X_test) # use the fit model to predict the c
2 y_test_pred_LR
```

```
Out[18]: array([16.5653963 , 21.18822792, 21.55107058, 10.88923816, 22.20231988,
13.35556872, 21.19692502, 7.35028523, 13.27547079, 15.12449511,
9.01443026, 6.52542825, 14.30205991, 8.97026042, 9.45679576,
12.00454351, 8.91549403, 16.15619251, 10.29582883, 18.72473553,
19.76821818, 13.77469028, 12.49638908, 21.53501762, 7.60860741,
5.6119801 , 20.91759483, 11.80627665, 9.08076637, 8.51412012,
12.17604891, 9.9691939 , 21.73008956, 12.77770578, 18.1011362 ,
20.07590796, 14.26202556, 20.93826535, 10.83938827, 4.38190607,
9.51332406, 12.40486324, 10.17045434, 8.09081363, 13.16388427,
5.2243552 , 9.28893833, 14.09330719, 8.69024497, 11.66119763,
15.71848432, 11.63156862, 13.35360735, 11.1531472 , 6.33636845,
9.76157954, 9.4195714 , 24.25516546, 7.69519137, 12.15317572])
```

```
In [19]: 1 # compare the y_test data and y_test_pred_LR data to check the performa
2
3 R2=r2_score(y_test,y_test_pred_LR)
4 MSE=mean_absolute_error(y_test,y_test_pred_LR)
5 RMSE=np.sqrt(MSE)
6 print('R-Square: ',R2)
7 print('Mean Square Error: ',MSE)
8 print('Root Mean Square Error: ',RMSE)
```

R-Square: 0.8609466508230368
Mean Square Error: 1.5116692224549084
Root Mean Square Error: 1.229499582128806

```
In [20]: 1 LR.coef_ # this gives the coefficient of each independent variable in t
```

```
Out[20]: array([0.04405928, 0.1992875 , 0.00688245])
```

```
In [21]: 1 LR.intercept_ # this gives the intercept of the linear regression equat
```

```
Out[21]: 2.70894909251591
```

Sales = 2.70894909251591 + 0.04405928 * TV + 0.1992875 * Radio + 0.00688245 * Newspaper

B) Decision Tree

```
In [22]: 1 # using a linear regression model
2
3 from sklearn.tree import DecisionTreeRegressor
4 DT=DecisionTreeRegressor()
5 DT.fit(X_train,y_train) # fit the train data into the model
```

Out[22]: DecisionTreeRegressor()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [23]: 1 y_train_pred_DT=DT.predict(X_train) # use the fit model to predict the
2 y_train_pred_DT
```

Out[23]: array([15. , 15.5, 11.9, 9.2, 12.8, 6.6, 20.7, 6.7, 11.2, 9.5, 10.7,
 8.8, 14.9, 17.1, 15.9, 15.9, 15. , 20.7, 11.7, 21.8, 9.3, 14.2,
 10.6, 17.3, 9.9, 15.2, 13.3, 25.4, 10.8, 24.2, 5.6, 19.2, 23.8,
 17.4, 17.4, 17.3, 12.9, 11.8, 15.9, 15.5, 12.9, 9.6, 20.1, 10.3,
 19.8, 11.4, 8.6, 15.7, 14.1, 13.2, 8. , 22.1, 25.5, 19. , 18.3,
 15.2, 18. , 8.5, 18. , 18.9, 16.6, 5.3, 3.2, 15.3, 12. , 15.5,
 17.6, 25.4, 12.3, 21.5, 17.1, 23.2, 11.5, 15.6, 12.9, 9.6, 18.5,
 19.2, 21.2, 19.6, 10.5, 11.8, 22.6, 9.7, 11.8, 9.5, 13.2, 13.4,
 7.3, 13.6, 20.2, 12.2, 16.7, 4.8, 9.7, 18.4, 14.5, 12.7, 24.4,
 8.8, 27. , 10.8, 12.2, 26.2, 14. , 20.2, 8.7, 11.4, 9.7, 23.8,
 14.8, 12.9, 12.5, 13.2, 14.4, 5.9, 9.7, 14.7, 10.1, 19.6, 10.4,
 22.6, 10.1, 1.6, 11.6, 14.7, 17.2, 12.2, 16. , 17. , 7. , 13.4,
 18. , 15.9, 12.4, 7.2, 19. , 19.4, 12.6, 14.8])

```
In [25]: 1 y_train
```

Out[25]: 169 15.0
97 15.5
31 11.9
12 9.2
35 12.8
...
106 7.2
14 19.0
92 19.4
179 12.6
102 14.8
Name: Sales, Length: 140, dtype: float64

```
In [24]: 1 # compare the y_train data and y_train_pred_LR data to check the perfor
2
3 R2=r2_score(y_train,y_train_pred_DT)
4 MSE=mean_absolute_error(y_train,y_train_pred_DT)
5 RMSE=np.sqrt(MSE)
6 print('R-Square: ',R2)
7 print('Mean Square Error: ',MSE)
8 print('Root Mean Square Error: ',RMSE)
```

R-Square: 1.0
Mean Square Error: 0.0
Root Mean Square Error: 0.0

```
In [26]: 1 y_test_pred_DT=DT.predict(X_test) # use the fit model to predict the ou
        2 y_test_pred_DT
```

```
Out[26]: array([18.5, 23.8, 17.6,  5.3, 23.8, 15.3, 22.6, 10.1, 12. , 16.6,  8.8,
                8.6, 11.7,  3.2, 10.1, 12.9,  5.3, 17.3,  9.7, 20.2, 17.6, 15.3,
                10.8, 23.8,  9.9,  8.7, 22.6, 12.2, 10.6,  4.8, 11.4,  9.9, 23.8,
                9.3, 15.9, 23.8, 10.4, 17.6, 12.4,  6.7, 11.2, 12.6, 10.4,  9.9,
                11.8,  8.8,  9.7, 15. ,  9.7, 11.8, 13.6, 12.9,  5.3,  5.3,  8.8,
                11.6,  9.9, 24.4,  6.6, 11.4])
```

```
In [27]: 1 # compare the y_test data and y_test_pred_LR data to check the performa
        2
        3 R2=r2_score(y_test,y_test_pred_DT)
        4 MSE=mean_absolute_error(y_test,y_test_pred_DT)
        5 RMSE=np.sqrt(MSE)
        6 print('R-Square: ',R2)
        7 print('Mean Square Error: ',MSE)
        8 print('Root Mean Square Error: ',RMSE)
```

R-Square: 0.9352062975938433

Mean Square Error: 1.0216666666666672

Root Mean Square Error: 1.010775280003754

C) Random Forest

```
In [31]: 1 # using a random forest model
        2
        3 from sklearn.ensemble import RandomForestRegressor
        4 RF=RandomForestRegressor()
        5 RF.fit(X_train,y_train)
```

```
Out[31]: RandomForestRegressor()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [32]: 1 y_train_pred_RF=RF.predict(X_train) # use the fit model to predict the
2 y_train_pred_RF
```

```
Out[32]: array([14.893, 15.507, 12.28 ,  9.131, 12.744,  6.719, 20.347,  6.714,
11.331,  9.151, 10.694,  8.901, 14.967, 17.591, 15.983, 15.799,
15.519, 20.476, 11.996, 22.171,  8.477, 14.445, 10.429, 17.486,
 9.97 , 14.979, 13.094, 25.347, 10.692, 24.491,  5.657, 18.993,
23.216, 16.821, 17.272, 17.145, 13.032, 11.558, 15.742, 15.47 ,
13.551,  9.618, 19.711, 10.457, 19.372, 11.273,  9.02 , 15.498,
14.275, 13.192,  7.949, 22.047, 25.237, 18.941, 18.278, 15.194,
17.46 ,  8.26 , 18.108, 19.158, 16.687,  5.442,  3.496, 14.772,
12.162, 15.474, 17.122, 24.666, 12.264, 21.696, 17.093, 23.07 ,
11.849, 15.62 , 12.858,  9.835, 18.307, 18.758, 20.797, 19.592,
10.93 , 11.886, 22.289,  9.776, 11.717,  9.462, 13.054, 13.049,
 7.729, 13.384, 20.089, 12.203, 16.552,  4.515,  9.819, 18.322,
14.712, 12.363, 24.126,  9.012, 26.413, 11.174, 12.187, 25.901,
13.295, 20.22 ,  9.007, 11.74 ,  9.62 , 23.223, 15.028, 12.351,
12.387, 13.389, 14.312,  6.218,  9.848, 14.427, 10.376, 19.34 ,
10.667, 22.156, 10.062,  3.983, 11.723, 14.824, 17.465, 12.501,
15.896, 17.114,  7.232, 13.28 , 18.413, 15.555, 12.233,  7.524,
18.83 , 19.381, 12.931, 14.437])
```

```
In [33]: 1 # compare the y_train data and y_train_pred_LR data to check the perfor
2
3 R2=r2_score(y_train,y_train_pred_RF)
4 MSE=mean_absolute_error(y_train,y_train_pred_RF)
5 RMSE=np.sqrt(MSE)
6 print('R-Square: ',R2)
7 print('Mean Square Error: ',MSE)
8 print('Root Mean Square Error: ',RMSE)
```

```
R-Square:  0.9954046126694128
Mean Square Error:  0.23939285714285685
Root Mean Square Error:  0.48927789357670437
```

```
In [34]: 1 y_test_pred_RF=RF.predict(X_test) # use the fit model to predict the ou
2 y_test_pred_RF
```

```
Out[34]: array([17.455, 22.02 , 20.401,  6.315, 23.268, 13.091, 22.686,  9.715,
12.361, 16.106,  8.421,  9.056, 12.191,  4.562, 10.532, 12.465,
 5.816, 16.739, 11.095, 19.596, 19.944, 12.988, 10.614, 22.108,
 9.97 ,  8.994, 22.32 , 12.536, 10.18 ,  5.126, 11.558, 10.838,
22.384,  8.477, 15.199, 20.349, 11.954, 20.515, 12.416,  7.585,
11.574, 13.004, 10.357,  9.698, 11.996,  9.013, 10.726, 15.508,
10.796, 11.842, 13.651, 12.604,  6.404,  5.936,  9.002, 11.409,
10.763, 25.265,  6.827, 12.227])
```

```
In [35]: 1 # compare the y_train data and y_train_pred_LR data to check the perfor
2
3 R2=r2_score(y_test,y_test_pred_RF)
4 MSE=mean_absolute_error(y_test,y_test_pred_RF)
5 RMSE=np.sqrt(MSE)
6 print('R-Square: ',R2)
7 print('Mean Square Error: ',MSE)
8 print('Root Mean Square Error: ',RMSE)
```

R-Square: 0.9830724032237538

Mean Square Error: 0.5809333333333325

Root Mean Square Error: 0.7621898276238883

```
In [ ]: 1
```